
Chapter 5: Overview of Query Processing

- Query Processing Overview
- Query Optimization
- Distributed Query Processing Steps

Acknowledgements: I am indebted to Arturas Mazeika for providing me his slides of this course.

Query Processing Overview

- **Query processing:** A 3-step process that transforms a high-level query (of relational calculus/SQL) into an **equivalent** and **more efficient** lower-level query (of relational algebra).

1. Parsing and translation

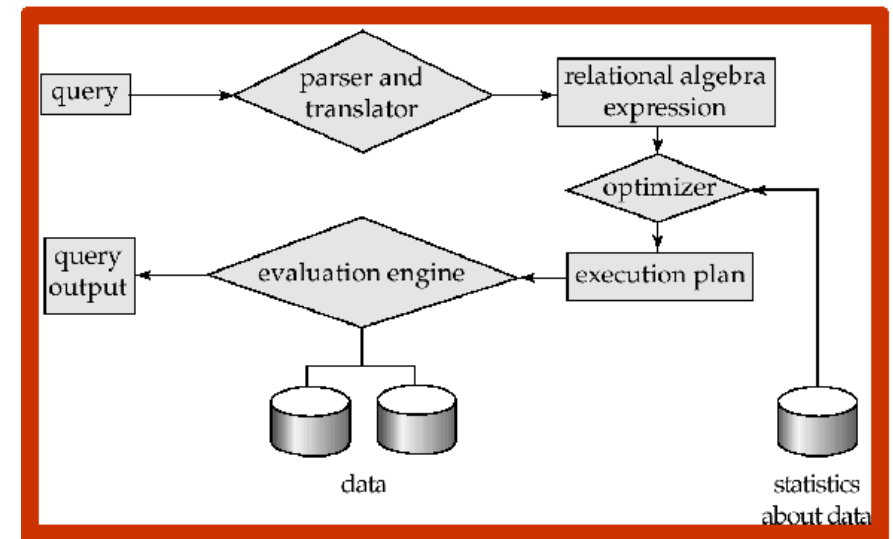
- Check syntax and verify relations.
- Translate the query into an equivalent relational algebra expression.

2. Optimization

- Generate an optimal evaluation plan (with lowest cost) for the query plan.

3. Evaluation

- The query-execution engine takes an (optimal) evaluation plan, executes that plan, and returns the answers to the query.



- The success of RDBMSs is due, in part, to the availability
 - of declarative query languages that allow to easily express complex queries without knowing about the details of the physical data organization and
 - of advanced query processing technology that transforms the high-level user/application queries into efficient lower-level query execution strategies.
- The query transformation should achieve both **correctness** and **efficiency**
 - The main difficulty is to achieve the efficiency
 - This is also one of the most important tasks of any DBMS
- **Distributed query processing:** Transform a high-level query (of relational calculus/SQL) on a distributed database (i.e., a set of global relations) into an **equivalent** and **efficient** lower-level query (of relational algebra) on relation fragments.
- Distributed query processing is more complex
 - Fragmentation/replication of relations
 - Additional communication costs
 - Parallel execution

Query Processing Example

- **Example:** Transformation of an SQL-query into an RA-query.

Relations: EMP(ENO, ENAME, TITLE), ASG(ENO,PNO,RESP,DUR)

Query: *Find the names of employees who are managing a project?*

- High level query

```
SELECT  ENAME
FROM    EMP , ASG
WHERE    EMP.ENO = ASG.ENO AND DUR > 37
```

- Two possible transformations of the query are:

- * Expression 1: $\Pi_{ENAME}(\sigma_{DUR>37 \wedge EMP.ENO=ASG.ENO}(EMP \times ASG))$

- * Expression 2: $\Pi_{ENAME}(EMP \bowtie_{ENO} (\sigma_{DUR>37}(ASG)))$

- Expression 2 avoids the expensive and large intermediate Cartesian product, and therefore typically is better.

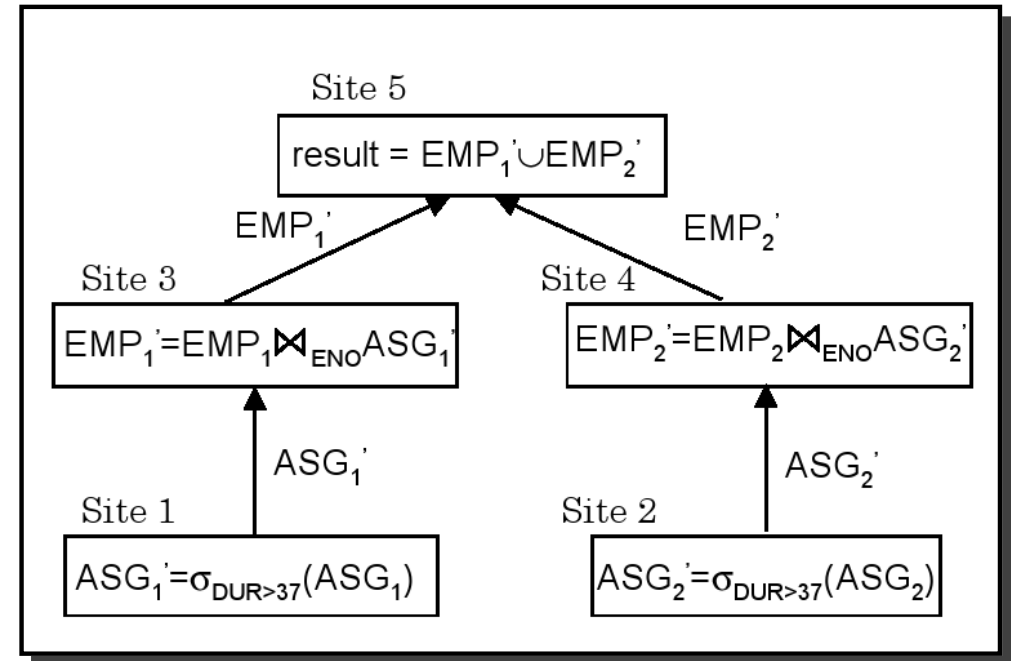
- We make the following assumptions about the data fragmentation
 - Data is (horizontally) fragmented:
 - * Site1: $ASG1 = \sigma_{ENO \leq "E3"}(ASG)$
 - * Site2: $ASG2 = \sigma_{ENO > "E3"}(ASG)$
 - * Site3: $EMP1 = \sigma_{ENO \leq "E3"}(EMP)$
 - * Site4: $EMP2 = \sigma_{ENO > "E3"}(EMP)$
 - * Site5: Result
 - Relations ASG and EMP are fragmented in the same way
 - Relations ASG and EMP are locally clustered on attributes RESP and ENO, respectively

Query Processing Example ...

- Now consider the expression $\Pi_{ENAME}(EMP \bowtie_{ENO} (\sigma_{DUR>37}(ASG)))$

- Strategy 1 (partially parallel execution):

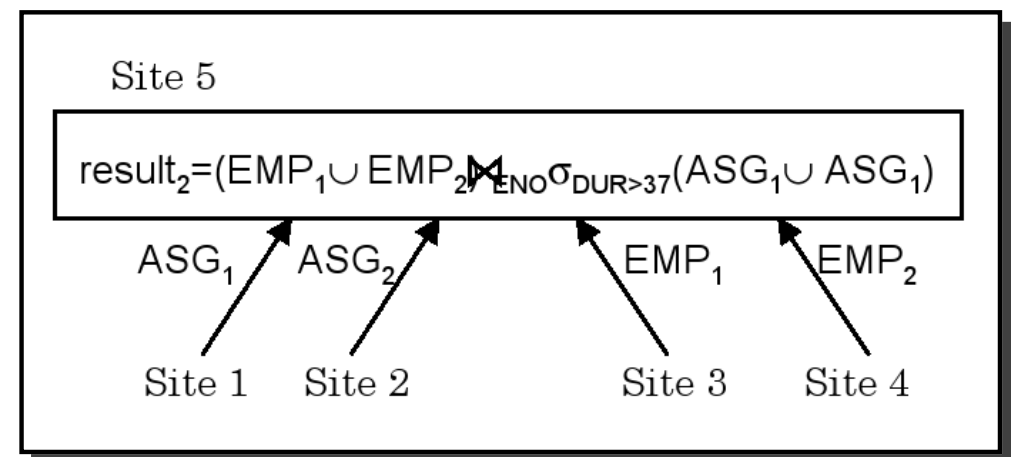
- Produce ASG'_1 and move to Site 3
- Produce ASG'_2 and move to Site 4
- Join ASG'_1 with EMP_1 at Site 3 and move the result to Site 5
- Join ASG'_2 with EMP_2 at Site 4 and move the result to Site 5
- Union the result in Site 5



- Strategy 2:

- Move ASG_1 and ASG_2 to Site 5
- Move EMP_1 and EMP_2 to Site 5
- Select and join at Site 5

- For simplicity, the final projection is omitted.



Query Processing Example ...

- Calculate the cost of the two strategies under the following assumptions:
 - Tuples are uniformly distributed to the fragments; 20 tuples satisfy $DUR > 37$
 - $size(EMP) = 400$, $size(ASG) = 1000$
 - tuple access cost = 1 unit; tuple transfer cost = 10 units
 - ASG and EMP have a local index on DUR and ENO
- Strategy 1
 - Produce ASG's: $(10+10) * \text{tuple access cost}$ 20
 - Transfer ASG's to the sites of EMPs: $(10+10) * \text{tuple transfer cost}$ 200
 - Produce EMP's: $(10+10) * \text{tuple access cost} * 2$ 40
 - Transfer EMP's to result site: $(10+10) * \text{tuple transfer cost}$ 200
 - Total cost 460
- Strategy 2
 - Transfer EMP_1, EMP_2 to site 5: $400 * \text{tuple transfer cost}$ 4,000
 - Transfer ASG_1, ASG_2 to site 5: $1000 * \text{tuple transfer cost}$ 10,000
 - Select tuples from $ASG_1 \cup ASG_2$: $1000 * \text{tuple access cost}$ 1,000
 - Join EMP and ASG': $400 * 20 * \text{tuple access cost}$ 8,000
 - Total cost 23,000

Query Optimization

- **Query optimization** is a crucial and difficult part of the overall query processing
- Objective of query optimization is to **minimize** the following cost function:

$$\text{I/O cost} + \text{CPU cost} + \text{communication cost}$$

- Two different scenarios are considered:
 - Wide area networks
 - * Communication cost dominates
 - low bandwidth
 - low speed
 - high protocol overhead
 - * Most algorithms ignore all other cost components
 - Local area networks
 - * Communication cost not that dominant
 - * Total cost function should be considered

Query Optimization ...

- **Ordering of the operators** of relational algebra is crucial for efficient query processing
- Rule of thumb: move expensive operators at the end of query processing
- Cost of RA operations:

Operation	Complexity
Select, Project (without duplicate elimination)	$O(n)$
Project (with duplicate elimination)	$O(n \log n)$
Group Join Semi-join Division Set Operators	$O(n \log n)$
Cartesian Product	$O(n^2)$

Several issues have to be considered in query optimization

- Types of query optimizers
 - wrt the search techniques (exhaustive search, heuristics)
 - wrt the time when the query is optimized (static, dynamic)
- Statistics
- Decision sites
- Network topology
- Use of semijoins

- **Types of Query Optimizers wrt Search Techniques**
 - Exhaustive search
 - * Cost-based
 - * Optimal
 - * Combinatorial complexity in the number of relations
 - Heuristics
 - * Not optimal
 - * Regroups common sub-expressions
 - * Performs selection, projection first
 - * Replaces a join by a series of semijoins
 - * Reorders operations to reduce intermediate relation size
 - * Optimizes individual operations

- **Types of Query Optimizers wrt Optimization Timing**

- Static

- * Query is optimized prior to the execution
 - * As a consequence it is difficult to estimate the size of the intermediate results
 - * Typically amortizes over many executions

- Dynamic

- * Optimization is done at run time
 - * Provides exact information on the intermediate relation sizes
 - * Have to re-optimize for multiple executions

- Hybrid

- * First, the query is compiled using a static algorithm
 - * Then, if the error in estimate sizes greater than threshold, the query is re-optimized at run time

- **Statistics**

- Relation/fragments
 - * Cardinality
 - * Size of a tuple
 - * Fraction of tuples participating in a join with another relation/fragment
- Attribute
 - * Cardinality of domain
 - * Actual number of distinct values
 - * Distribution of attribute values (e.g., histograms)
- Common assumptions
 - * Independence between different attribute values
 - * Uniform distribution of attribute values within their domain

- **Decision sites**

- Centralized

- * Single site determines the "best" schedule
 - * Simple
 - * Knowledge about the entire distributed database is needed

- Distributed

- * Cooperation among sites to determine the schedule
 - * Only local information is needed
 - * Cooperation comes with an overhead cost

- Hybrid

- * One site determines the global schedule
 - * Each site optimizes the local sub-queries

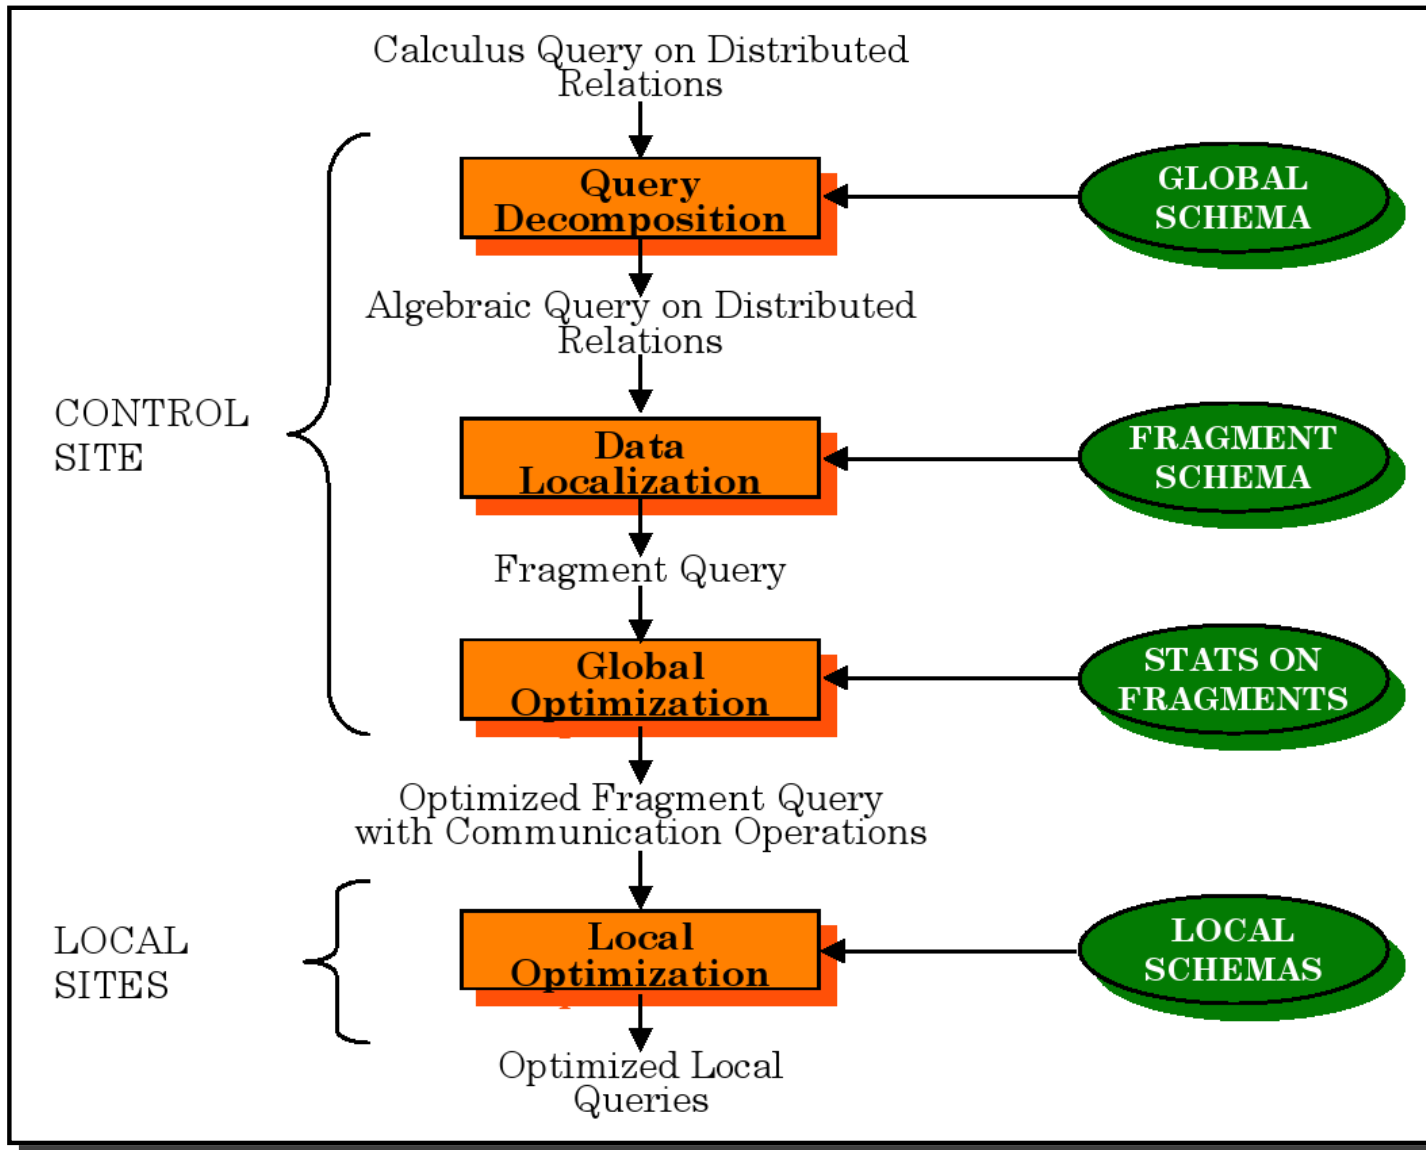
- **Network topology**

- Wide area networks (WAN) point-to-point
 - * Characteristics
 - Low bandwidth
 - Low speed
 - High protocol overhead
 - * Communication cost dominate; all other cost factors are ignored
 - * Global schedule to minimize communication cost
 - * Local schedules according to centralized query optimization
- Local area networks (LAN)
 - * Communication cost not that dominant
 - * Total cost function should be considered
 - * Broadcasting can be exploited (joins)
 - * Special algorithms exist for star networks

- **Use of Semijoins**

- Reduce the size of the join operands by first computing semijoins
- Particularly relevant when the main cost is the communication cost
- Improves the processing of distributed join operations by reducing the size of data exchange between sites
- However, the number of messages as well as local processing time is increased

Distributed Query Processing Steps



Conclusion

- Query processing transforms a high level query (relational calculus) into an equivalent lower level query (relational algebra). The main difficulty is to achieve the efficiency in the transformation
- Query optimization aims to minimize the cost function:

$$\text{I/O cost} + \text{CPU cost} + \text{communication cost}$$

- Query optimizers vary by search type (exhaustive search, heuristics) and by type of the algorithm (dynamic, static, hybrid). Different statistics are collected to support the query optimization process
- Query optimizers vary by decision sites (centralized, distributed, hybrid)
- Query processing is done in the following sequence: query decomposition → data localization → global optimization → local optimization